

1. **Set the range** to 2026-07-31 00:00 → 23:59 UTC – the unseen day. Outside it the panels are empty.
2. **Show the ACCURATE curve** – the ramp, and the peak at 11:17.
3. **Show NAIVE beside it** – same data, session-span reading. The gap is backgrounded and paused time.
4. **Apply a filter** – platform, or video_resolution. Watch the curve change shape, not just scale.
5. **Say the resolution point:** that column arrived with the unseen file. No migration. Filterable the same day.
6. **Show pipeline health** – our own ingestion and query latency, in ClickStack.

△ 11 of 12 filters move the curve. **Country does not** – it is the constant 'india' in all 7,000,000 rows. Say it before anyone asks.

△ **Latency on screen will not be 28 ms once you filter.** The unfiltered curve comes off the promoted delta tier at 28 ms. A filtered drilldown answers on arbitrary dimensions and takes ~4.4 s. Both are real and both are ours – say which path you are on rather than letting the number surprise you: "that is the generic path, it answers on a column we never pre-aggregated."

- the flexible slide

