

FeatureLens

Agentic product analytics that evolves instrumentation, semantic context, and trustworthy decisions together.

Plausible analytics is not the same as trustworthy analytics

THE GAP

A release changes the physical data contract and what the business can safely ask. Today those changes drift apart.
A fluent answer can still use the wrong event, grain, denominator, table, or context version.

Traditional analytics

Dashboards arrive after releases
Definitions drift across teams
Provenance is hard to audit

Generic AI analytics

Text-to-SQL is too permissive
Raw rows cross the trust boundary
Confidence can hide weak evidence

One release package. Three agents. No skipped gates.

01 · PHYSICAL

02 · SEMANTIC

03 · DECISION

Instrumentation Agent

Profiles arbitrary NDJSON, proposes typed ClickHouse DDL, validates it, then pauses for human approval.

Context Agent

Publishes an immutable graph only after schema, row count, event fingerprint, and blocking evolution checks pass.

Analytics Agent

Compiles ontology-linked aggregate SQL, produces a deterministic answer, and constrains optional LLM synthesis.

The product feels like a release workspace, not an agent demo

A PM-native control plane

Feature Releases, Decision Inbox, Ask FeatureLens, Context & Schemas, Pipeline Activity, and Trace Explorer share one governed backend. LibreChat can call the same seven MCP tools without owning business truth.



The system is already producing auditable evidence

LIVE PROOF

Six releases were replayed against retained ClickHouse data. Every run passed all nine evolution checks and exported a synced Langfuse trace.

34,982

retained feature events


54 / 54

context gates passed

4 / 4

standard PM probes traced

Business meaning grows with the product—and stays versioned



Context v0 → v6 preserves every parent node while adding verified feature semantics.

147

semantic nodes

396

explicit edges

The sealed feature generalized—and changed the product conclusion

PROMO / COUPON AT CHECKOUT

FeatureLens profiled 5,363 events, proposed a 19-column SharedMergeTree schema, inferred the success path, and kept coupon_rejected as a branch.

It then inferred a nullable cohort marker from the declared question and verified coupon_code field.

Coupon-marked cohort

848 entrants
366 completed checkout
43.16% completion

Null-marker baseline

1,252 entrants
621 completed checkout
49.60% completion

Trust is designed into every layer

ClickHouse is the evidence boundary

Only allowlisted aggregate SQL executes. Raw rows never enter the LLM prompt. Physical data and governance history remain inspectable together.

The LLM is outside the trust boundary

It rewrites a complete deterministic answer into concise PM language. Invalid output falls back without losing evidence or provenance.

Langfuse closes the feedback loop

Trace Explorer shows observations, cost, tokens, evaluations, and lets PMs write typed helpfulness and issue-category scores.

**Trust every answer that
follows.**

clickathon-2026.view26.com

6 releases · 54/54 gates

ClickHouse · Langfuse · LibreChat